



Apache Kylin Introduction

Dec 14, 2014

韩卿 | Luke Han

Co-creator of Apache Kylin | lukehan@apache.org

Sr. Product Manager, eBay CCOE



Agenda



- What's Apache Kylin?
- Feature & Tech Highlights
- Performance
- Open Source & Roadmap
- Q & A

What's Kylin



kylin / 'ki:'lin / 麒麟

--n. (in Chinese art) a mythical animal of composite form


Extreme OLAP Engine for Big Data

Kylin is an open source Distributed Analytics Engine from eBay that provides SQL interface and multi-dimensional analysis (OLAP) on Hadoop supporting extremely large datasets

- Open Sourced on Oct 1st, 2014
- Be accepted as Apache Incubator Project on Nov 25th, 2014



Big Data Era



- More and more data becoming available on Hadoop
- Limitations in existing Business Intelligence (BI) Tools
 - Limited support for Hadoop
 - Data size growing exponentially
 - High latency of interactive queries
 - Scale-Up architecture
- Challenges to adopt Hadoop as interactive analysis system
 - Majority of analyst groups are SQL savvy
 - No mature SQL interface on Hadoop
 - OLAP capability on Hadoop ecosystem not ready yet



Business Needs for Big Data Analysis

- Sub-second query latency on billions of rows
 - ANSI SQL for both analysts and engineers
 - Full OLAP capability to offer advanced functionality
 - Seamless Integration with BI Tools
-
- Support of high cardinality and high dimensions
 - High concurrency - thousands of end users
 - Distributed and scale out architecture for large data volume

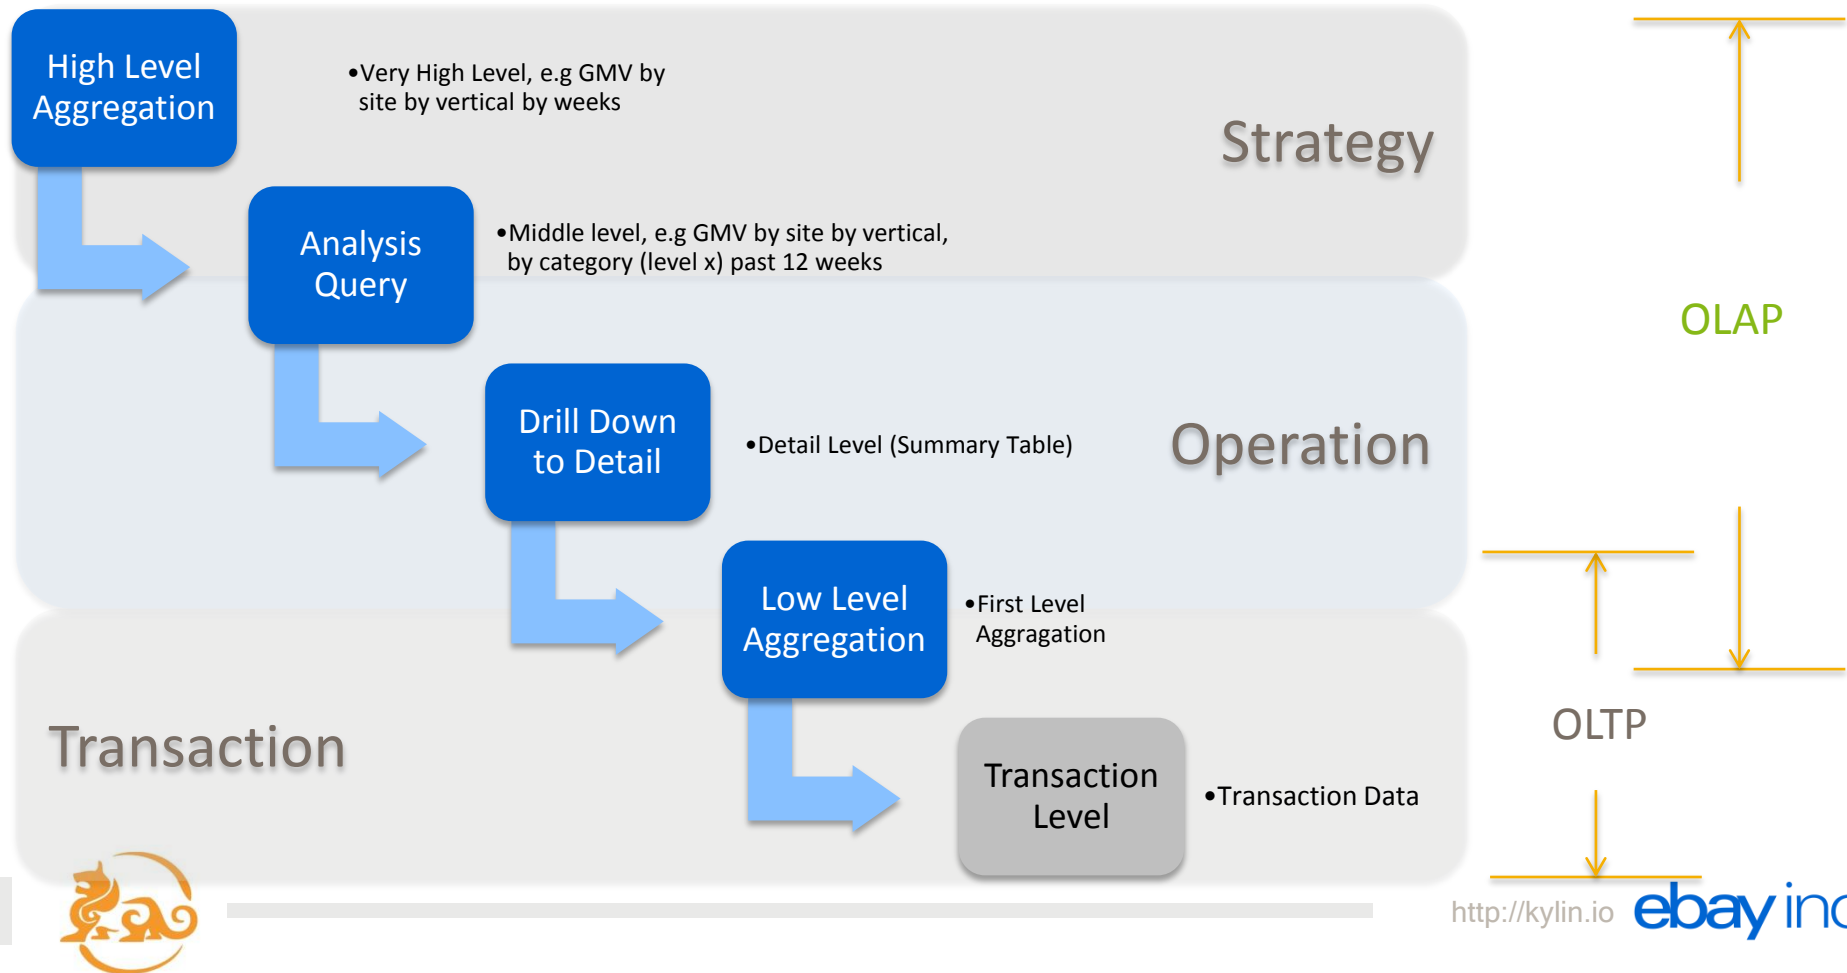




Why not
Build an engine from scratch?

Analytics Query Taxonomy

Kylin is designed to accelerate 80+% analytics queries performance on Hadoop




Technical Challenges

- Huge volume data
 - Table scan
- Big table joins
 - Data shuffling
- Analysis on different granularity
 - Runtime aggregation expensive
- Map Reduce job
 - Batch processing



OLAP Cube - Balance between Space and Time

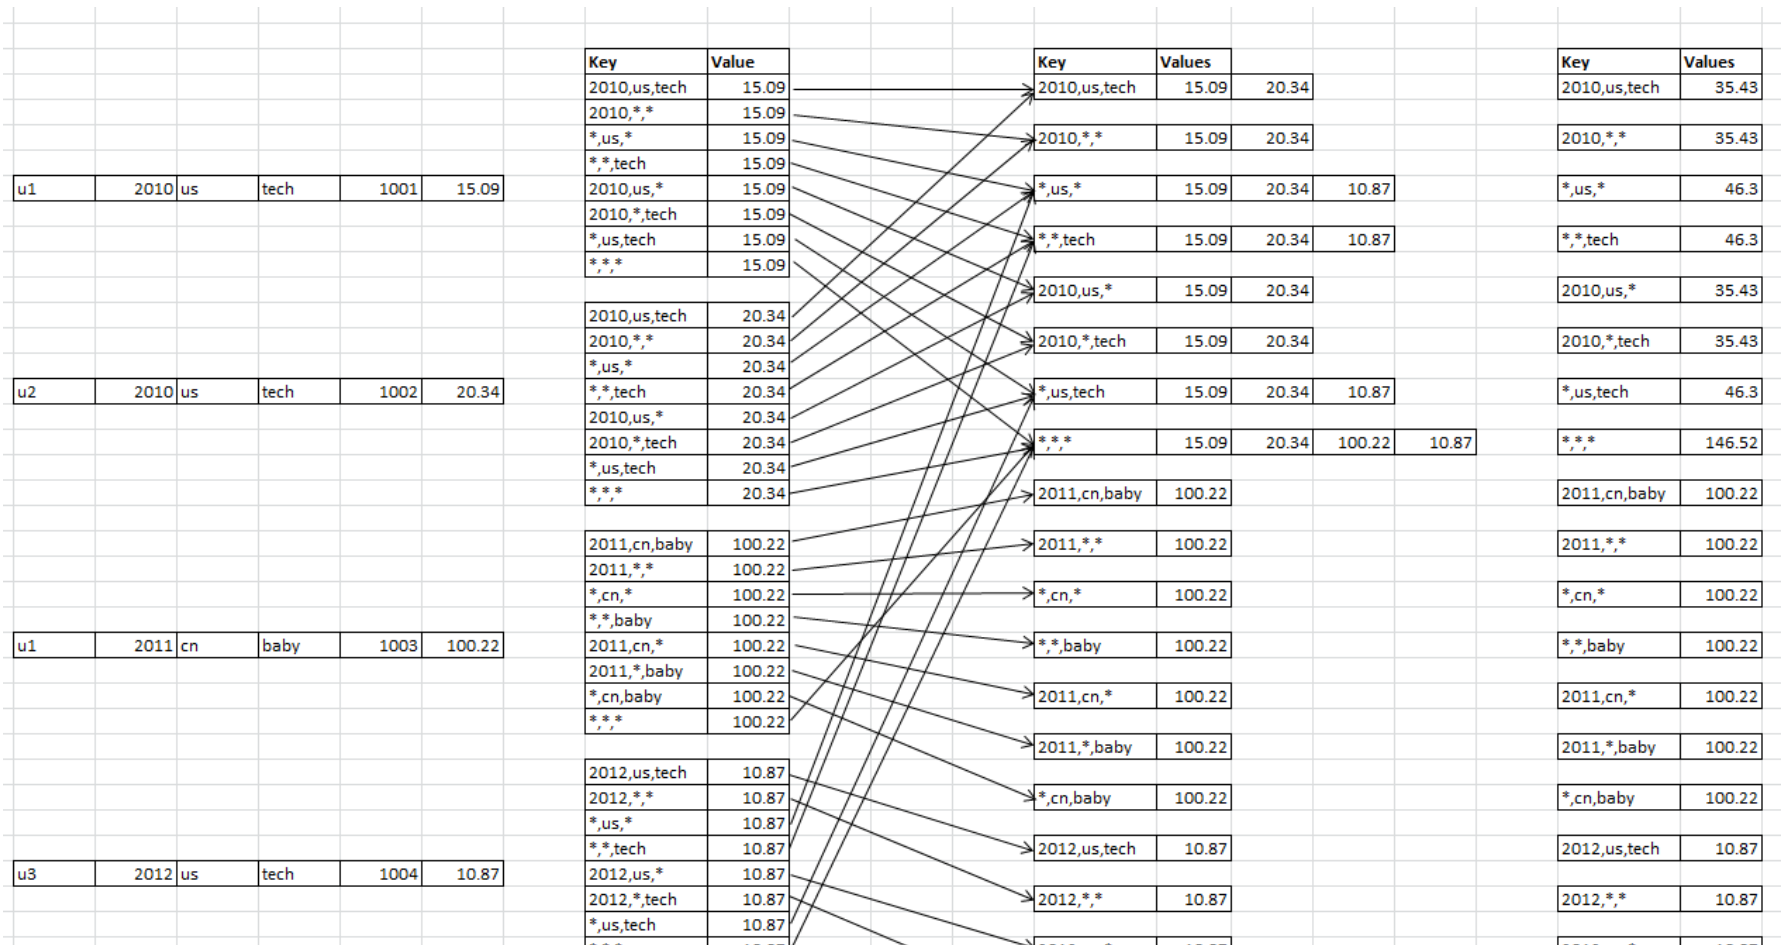
- Cuboid = one combination of dimensions
- Cube = all combination of dimensions (all cuboids)



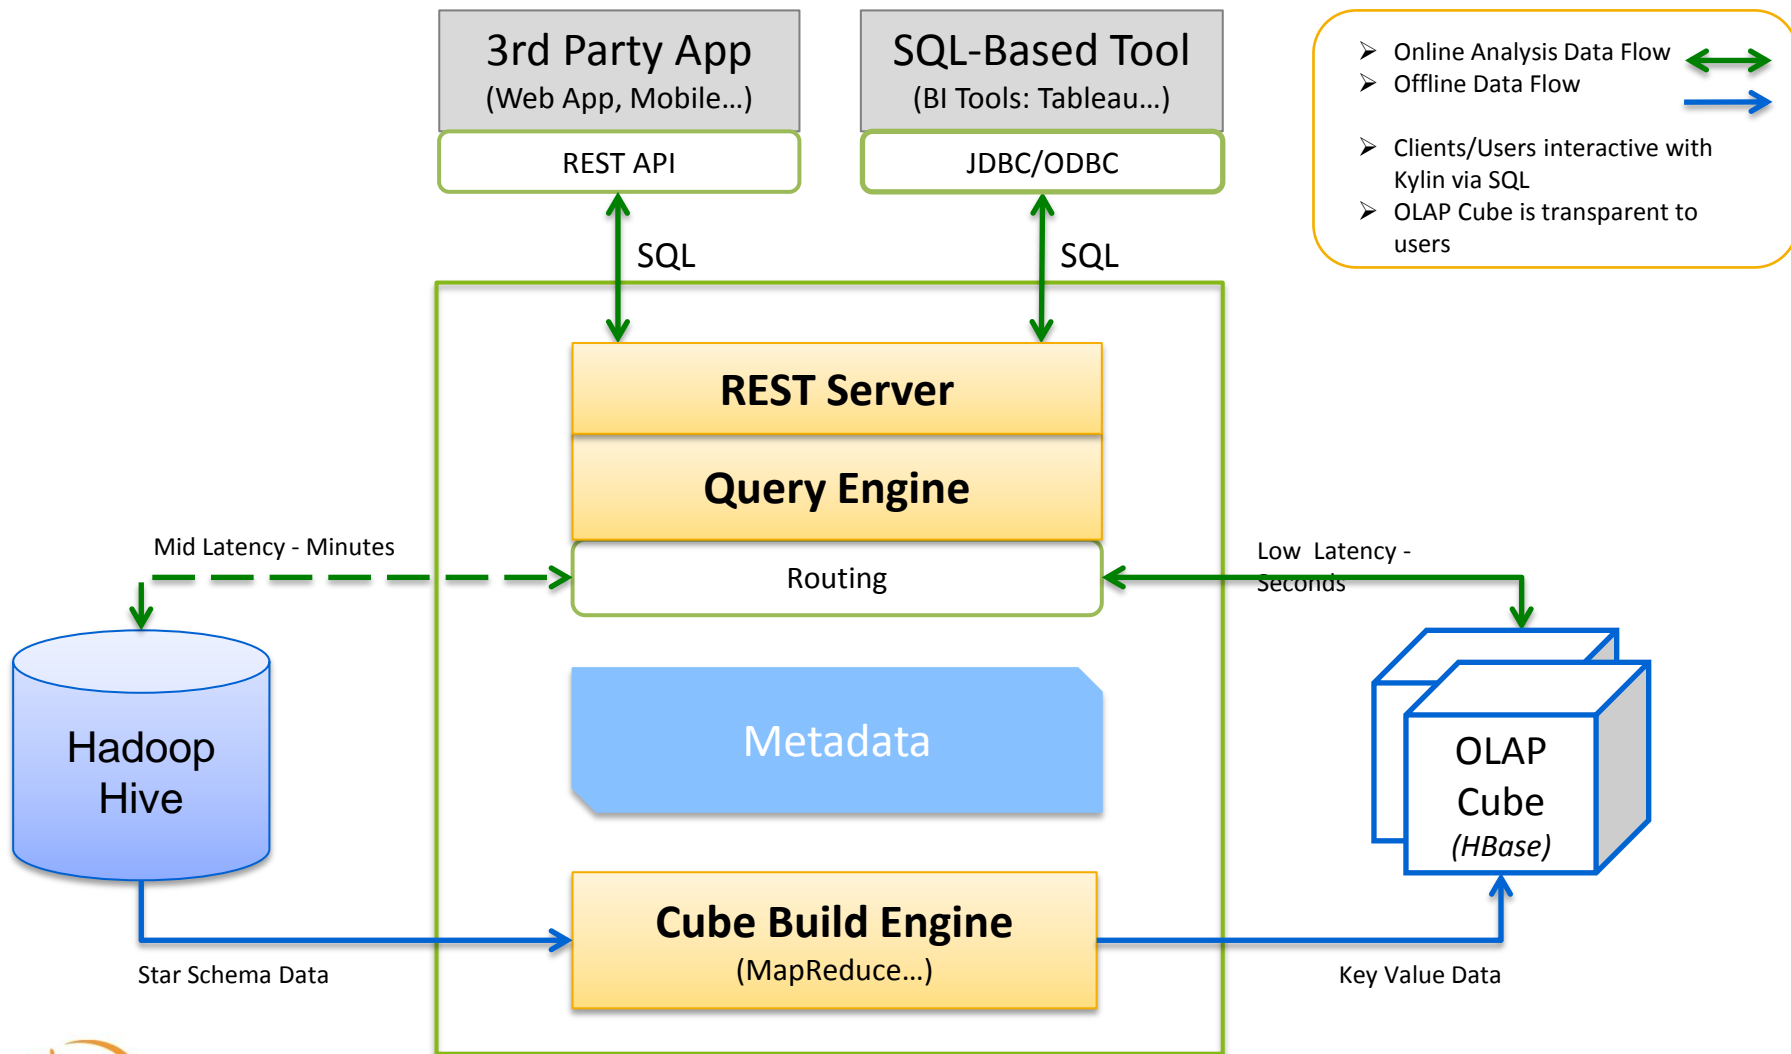
- Base vs. aggregate cells; ancestor vs. descendant cells; parent vs. child cells
 1. (9/15, milk, Urbana, Dairy_land) - **<time, item, location, supplier>**
 2. (9/15, milk, Urbana, *) - **<time, item, location>**
 3. (*, milk, Urbana, *) - **<item, location>**
 4. (*, milk, Chicago, *) - **<item, location>**
 5. (*, milk, *, *) - **<item>**



From Relational to Key-Value



Kylin Architecture Overview



How Does Kylin Utilize Hadoop Components?

- Hive
 - Input source
 - Pre-join star schema during cube building
- MapReduce
 - Pre-aggregation metrics during cube building
- HDFS
 - Store intermediated files during cube building.
- HBase
 - Store data cube.
 - Serve query on data cube.
 - Coprocessor is used for query processing.





Agenda



- What's Apache Kylin?
- Feature & Tech Highlights
- Performance
- Open Source & Roadmap
- Q & A

Features Highlights

- **Extremely Fast OLAP Engine at Scale**

Kylin is designed to reduce query latency on Hadoop for 10+ billions of rows of data

- **ANSI SQL Interface on Hadoop**

Kylin offers ANSI SQL on Hadoop and supports most ANSI SQL query functions

- **Seamless Integration with BI Tools**

Kylin currently offers integration capability with BI Tools like Tableau.

- **Interactive Query Capability**

Users can interact with Hadoop data via Kylin at sub-second latency, better than Hive queries for the same dataset

- **MOLAP Cube**

User can define a data model and pre-build in Kylin with more than 10+ billions of raw data records



Features Highlights...

- Compression and Encoding Support
- Incremental Refresh of Cubes
- Approximate Query Capability for distinct Count (HyperLogLog)
- Leverage HBase Coprocessor for query latency
- Job Management and Monitoring
- Easy Web interface to manage, build, monitor and query cubes
- Security capability to set ACL at Cube/Project Level
- Support LDAP Integration



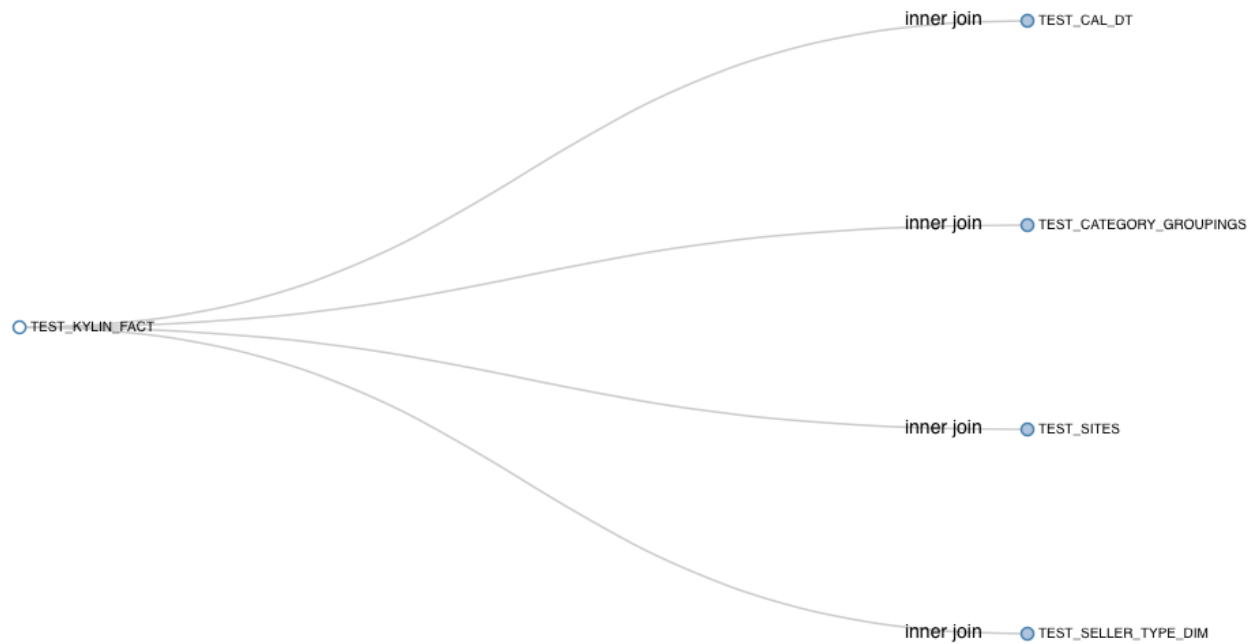
Cube Designer

Kylin Query Cubes Jobs Tables Admin Help Welcome, AD

Recent Tables Cube Designer

test_kylin_cube_with_slr_empty	READY	6.52 MB	10,000	2014-11-11 13:59:15	Action	Action
--------------------------------	-------	---------	--------	---------------------	--------	--------

Grid Visualization SQL JSON Access Notification HBase



Jobs	Tables	Admin	
------	--------	-------	--

	Sta
_empty	
empty	

```
SequenceID: 0
Status: FINISHED
Duration: 5.44 mins
Waiting: 0 seconds
Start At: 2014-11-0
End At: 2014-11-05
Data Size: 446.61 I
MR Job:
job_141516805639
```



Query and Visualization

Results (1215)

Visualization Export

Drag a column header here and drop it to group by that column.

LSTG_FORMAT...	WEEK_BEG_DT	META_CATEG...	PRICE
FP-GTC	2013-01-01	Sports MeCard...	93.268450871...
Auction	2013-01-01	Business & Ind...	55.635632631...
Others	2013-01-01	Sporting Goods	82.533676494...
Auction	2013-01-01	Real Estate	20.347844733...
Auction	2013-01-01	Toys & Hobbies	4.0906749147...


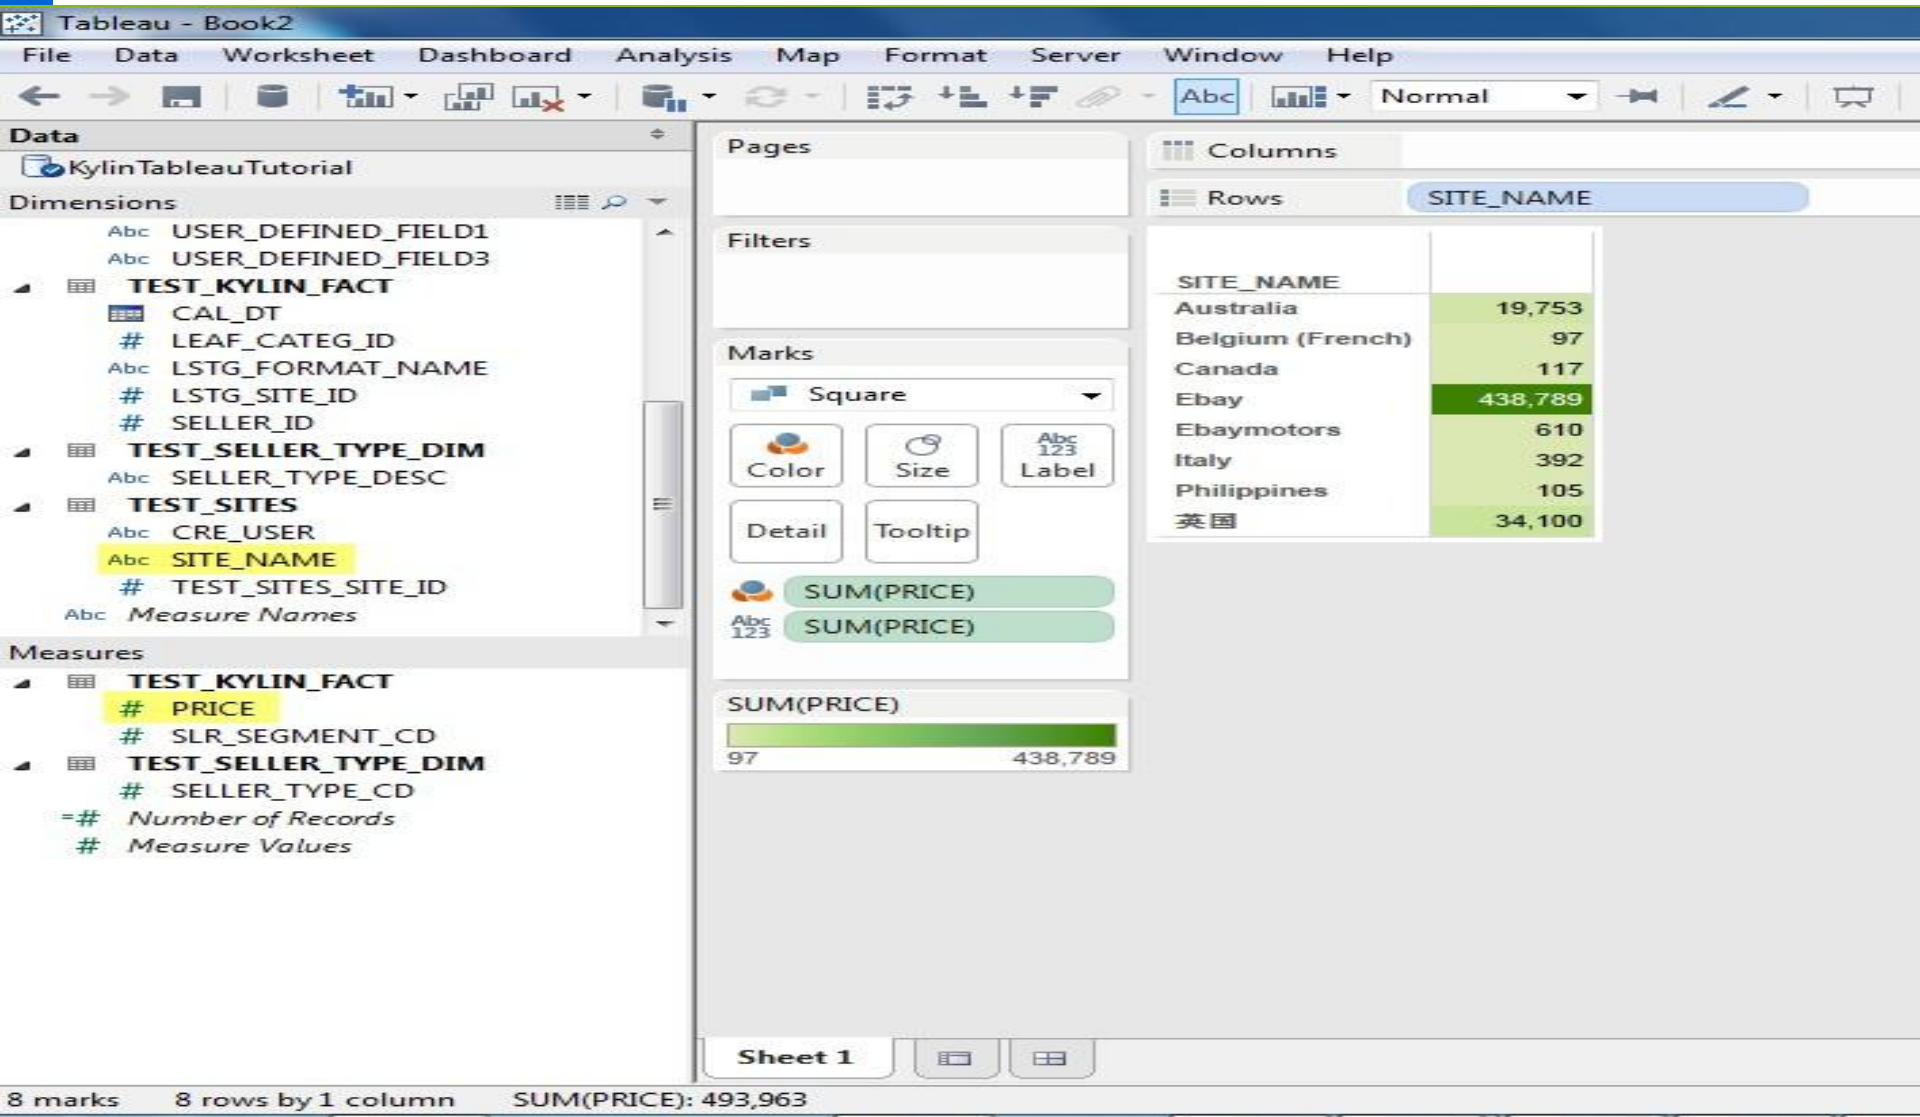



Tableau Integration



Data Modeling



End User



Source
Star Schema




Cube Modeler

Cube: ...
Fact Table: ...
Dimensions: ...
Measures: ...
Storage(HBase): ...

Mapping
Cube Metadata



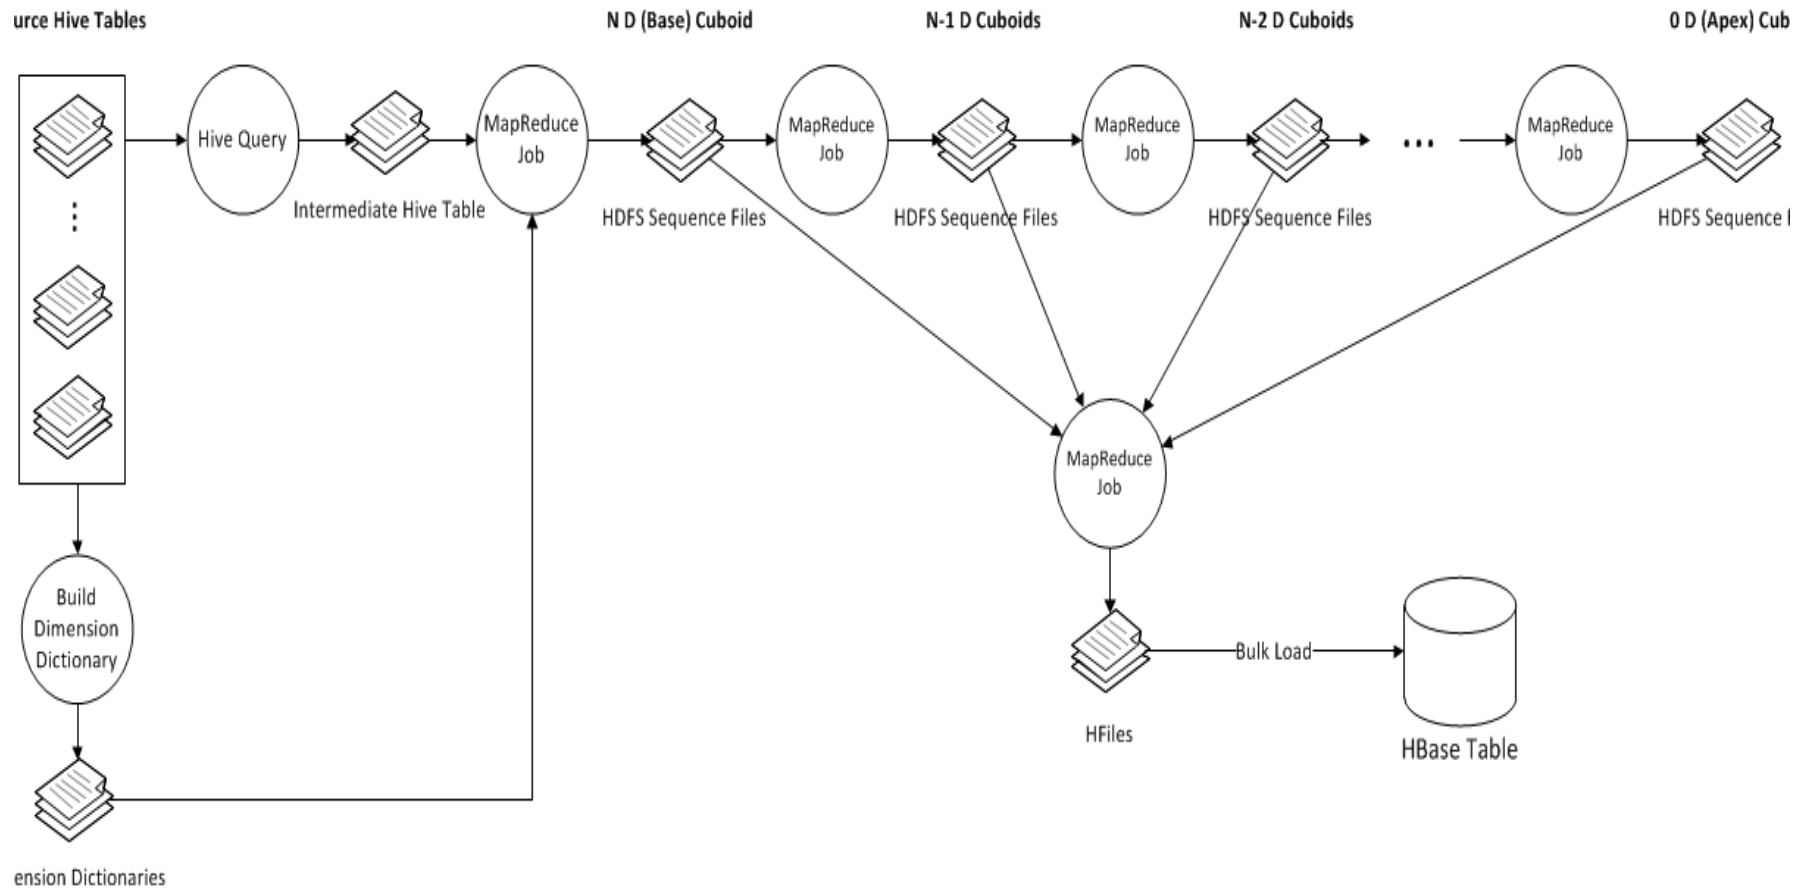
Admin



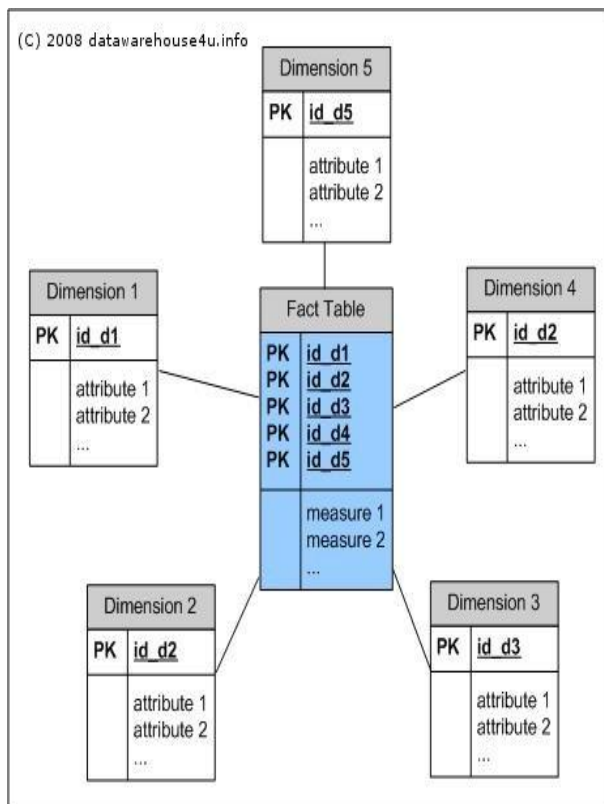
Target
HBase Storage



Cube Build Job Flow



How To Store Cube? - HBase Schema



Pre-Joined Flat Table	
Dimensions	D1
	D2
	...
	DX
Measures	M1
	M2
	...
	MY

Pre-Aggregated Table	
Dimensions	D3
	D6
	D8
Measures	M1
	M2
	...
	MY

Cuboid ID	D1	D2	...	DX
-----------	----	----	-----	----

M1	M2	...	MY
----	----	-----	----

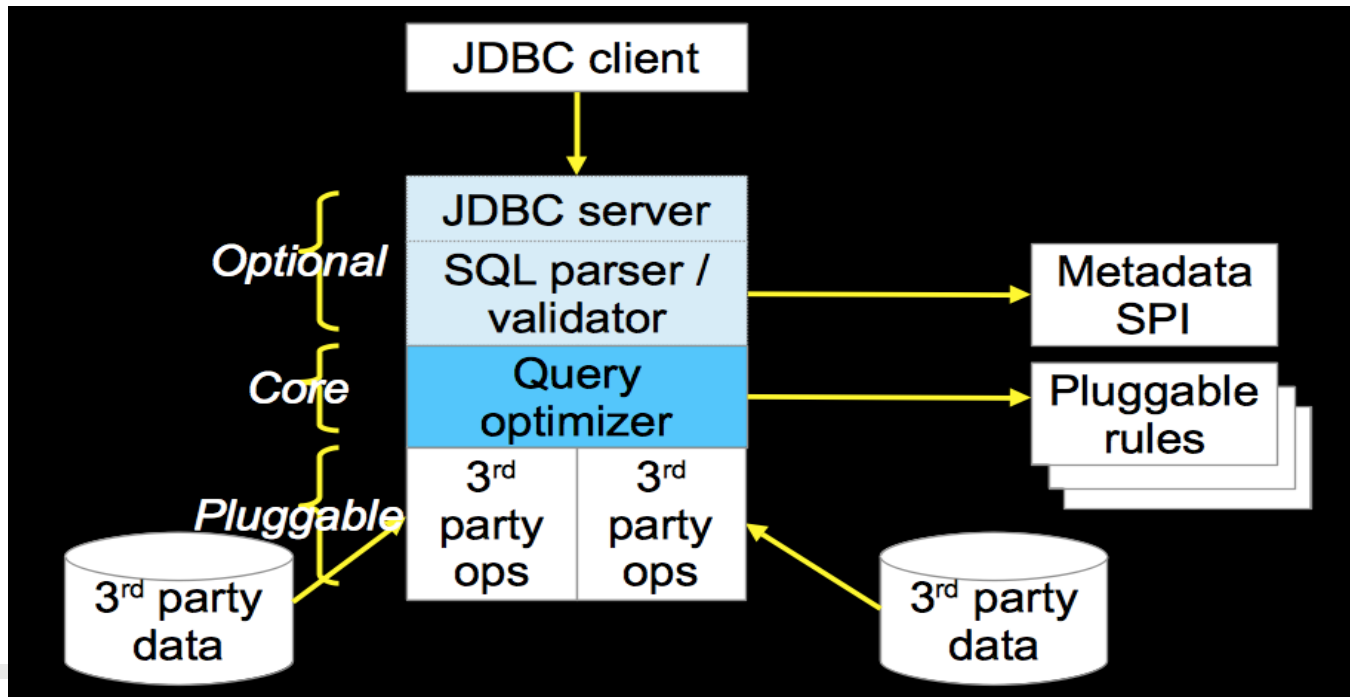
Row Key = Cuboid ID+Dimensions

Row Value = Measures



Query Engine - Calcite (Optiq)

- Dynamic data management framework.
- Formerly known as Optiq, Calcite is an Apache incubator project, used by Apache Drill and Apache Hive, among others.
- <http://optiq.incubator.apache.org>



Query Engine - Kylin Explain Plan

```
SELECT test_cal_dt.week_beg_dt, test_category.category_name, test_category.lvl2_name, test_category.lvl3_name,  
test_kylin_fact.lstg_format_name, test_sites.site_name, SUM(test_kylin_fact.price) AS GMV, COUNT(*) AS TRANS_CNT  
FROM test_kylin_fact  
  LEFT JOIN test_cal_dt ON test_kylin_fact.cal_dt = test_cal_dt.cal_dt  
  LEFT JOIN test_category ON test_kylin_fact.leaf_categ_id = test_category.leaf_categ_id AND test_kylin_fact.lstg_site_id =  
test_category.site_id  
  LEFT JOIN test_sites ON test_kylin_fact.lstg_site_id = test_sites.site_id  
WHERE test_kylin_fact.seller_id = 123456 OR test_kylin_fact.lstg_format_name = 'New'  
GROUP BY test_cal_dt.week_beg_dt, test_category.category_name, test_category.lvl2_name, test_category.lvl3_name,  
test_kylin_fact.lstg_format_name, test_sites.site_name
```

OLAPToEnumerableConverter

```
OLAPProjectRel(WEEK_BEG_DT=[0], category_name=[1], CATEG_LVL2_NAME=[2], CATEG_LVL3_NAME=[3],  
LSTG_FORMAT_NAME=[4], SITE_NAME=[5], GMV=[CASE(=(7, 0), null, 6)], TRANS_CNT=[8])  
  OLAPAggregateRel(group=[{0, 1, 2, 3, 4, 5}], agg#0=[SUM0(6)], agg#1=[COUNT(6)], TRANS_CNT=[COUNT()])  
    OLAPProjectRel(WEEK_BEG_DT=[13], category_name=[21], CATEG_LVL2_NAME=[15], CATEG_LVL3_NAME=[14],  
LSTG_FORMAT_NAME=[5], SITE_NAME=[23], PRICE=[0])  
      OLAPFilterRel(condition=[OR(=(3, 123456), =(5, 'New'))])  
        OLAPJoinRel(condition=[=(2, 25)], joinType=[left])  
          OLAPJoinRel(condition=[AND(=(6, 22), =(2, 17))], joinType=[left])  
            OLAPJoinRel(condition=[=(4, 12)], joinType=[left])  
              OLAPTableScan(table=[[DEFAULT, TEST_KYLIN_FACT]], fields=[[0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11]])  
                OLAPTableScan(table=[[DEFAULT, TEST_CAL_DT]], fields=[[0, 1]])  
                  OLAPTableScan(table=[[DEFAULT, test_category]], fields=[[0, 1, 2, 3, 4, 5, 6, 7, 8]])  
                    OLAPTableScan(table=[[DEFAULT, TEST_SITES]], fields=[[0, 1, 2]])
```



How To Optimize Cube? - Full Cube vs. Partial Cube

- **Full Cube**

- Pre-aggregate all dimension combinations
- “Curse of dimensionality”: N dimension cube has 2^N cuboid.

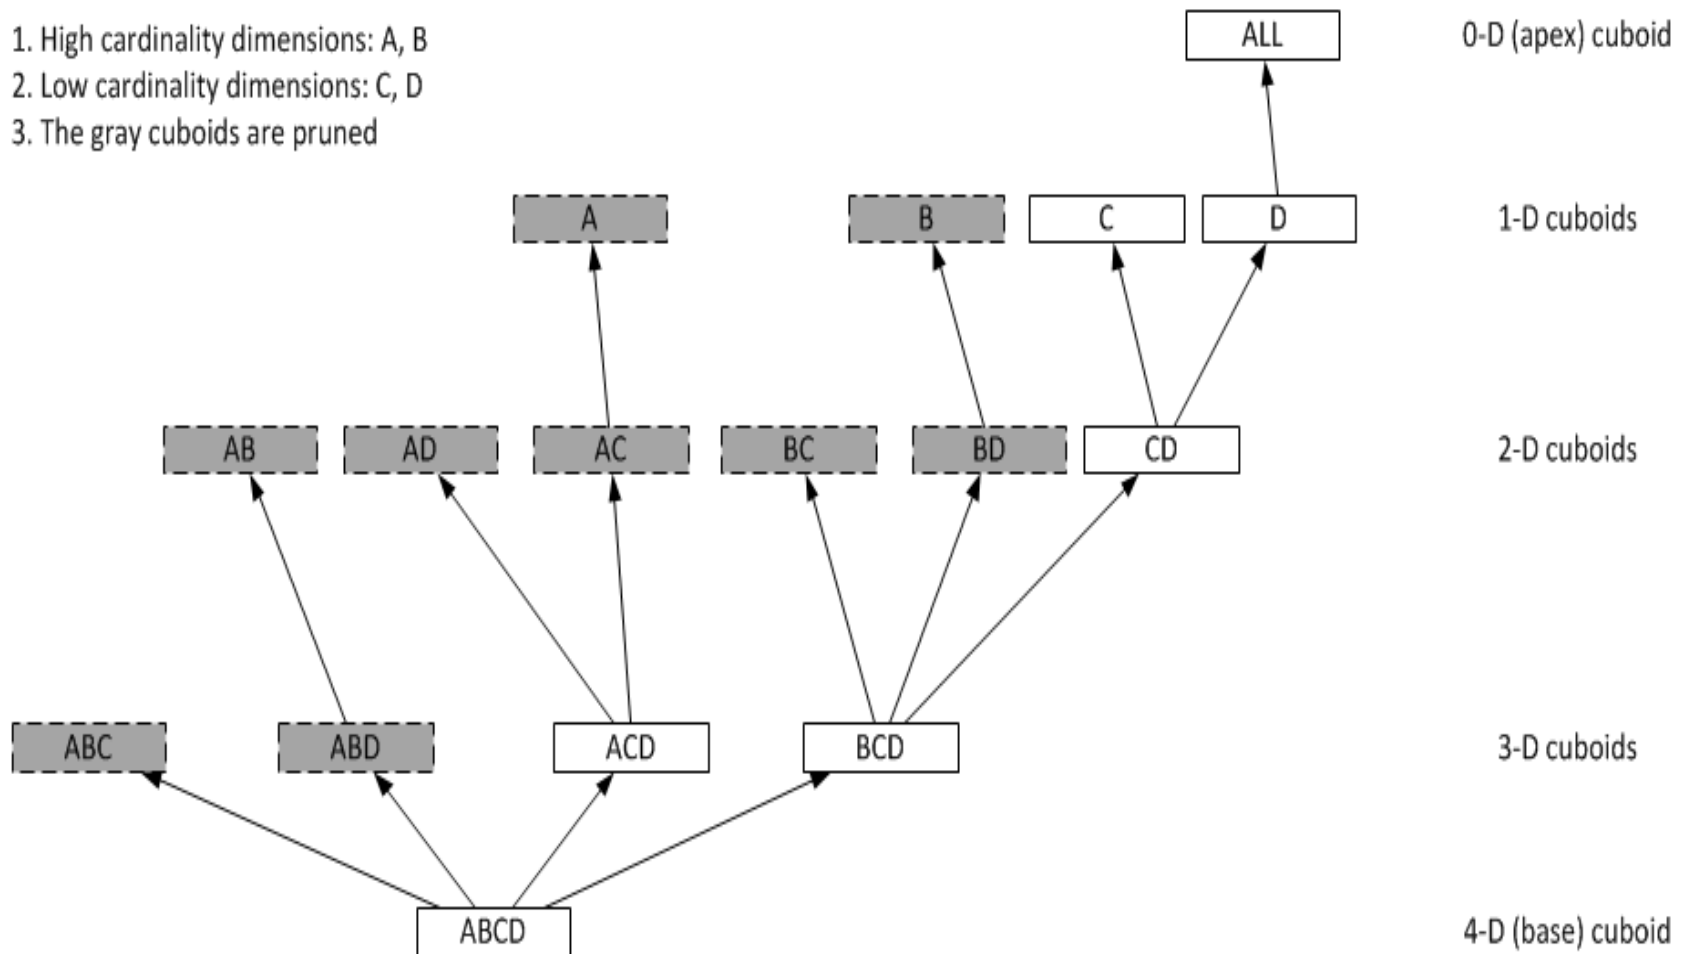
- **Partial Cube**

- To avoid dimension explosion, we divide the dimensions into different aggregation groups
 - $2^{N+M+L} \rightarrow 2^N + 2^M + 2^L$
- For cube with 30 dimensions, if we divide these dimensions into 3 group, the cuboid number will reduce from 1 Billion to 3 Thousands
 - $2^{30} \rightarrow 2^{10} + 2^{10} + 2^{10}$
- Tradeoff between online aggregation and offline pre-aggregation




How To Optimize Cube? - Partial Cube

1. High cardinality dimensions: A, B
2. Low cardinality dimensions: C, D
3. The gray cuboids are pruned




How To Optimize Cube? - Incremental Building


Aggregate the results when querying




cube-Y-2011:2012



cube-M-2013-1:8



cube-D-2013-09-1:20



cube-D-2013-09-21

1. Cube is immutable
2. Merge small cubes into a larger one



Inverted Index

- Challenge
 - Has no raw data records
 - Slow table scan on high cardinality dimensions
- Inverted Index Storage (an ongoing effort)
 - Persist the raw table
 - Bitmap inverted index
 - Time range partition
 - In-memory (block cache)
 - Parallel scan (endpoint coprocessor)




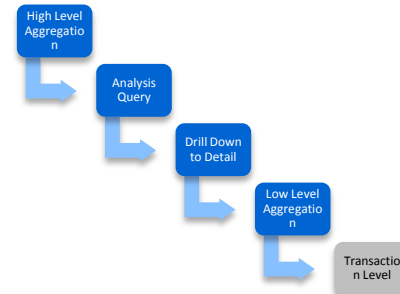


Agenda



- What's Apache Kylin?
- Feature & Tech Highlights
- Performance
- Open Source & Roadmap
- Q & A

Kylin vs. Hive





#	Query Type	Return Dataset	Query On Kylin (s)	Query On Hive (s)	Comments
1	High Level Aggregation	4	0.129	157.437	1,217 times
2	Analysis Query	22,669	1.615	109.206	68 times
3	Drill Down to Detail	325,029	12.058	113.123	9 times
4	Drill Down to Detail	524,780	22.42	6383.21	278 times
5	Data Dump	972,002	49.054	N/A	



Performance -- Concurrency

Single Tomcat Instance on a Single Machine

	Parallel Thread #	Data			Latency (ms)				Throughput
		Raw Recors	HBase Scan	Return	Min	Max	Median	90% Line	
High Level Aggregation Query	30	1,940,304,293	5	5	67	1809	334	355	72.5/sec
Detail Level Query (with Seller ID)	30	13,683,834,542	43934	7283	1758	4534	2182	3171	9.7/sec




Linear scale out with more nodes



Performance - Query Latency

90%tile queries <5s

Kylin Query Latency (90% and 95%)



Green Line: 90%tile queries

Gray Line: 95%tile queries





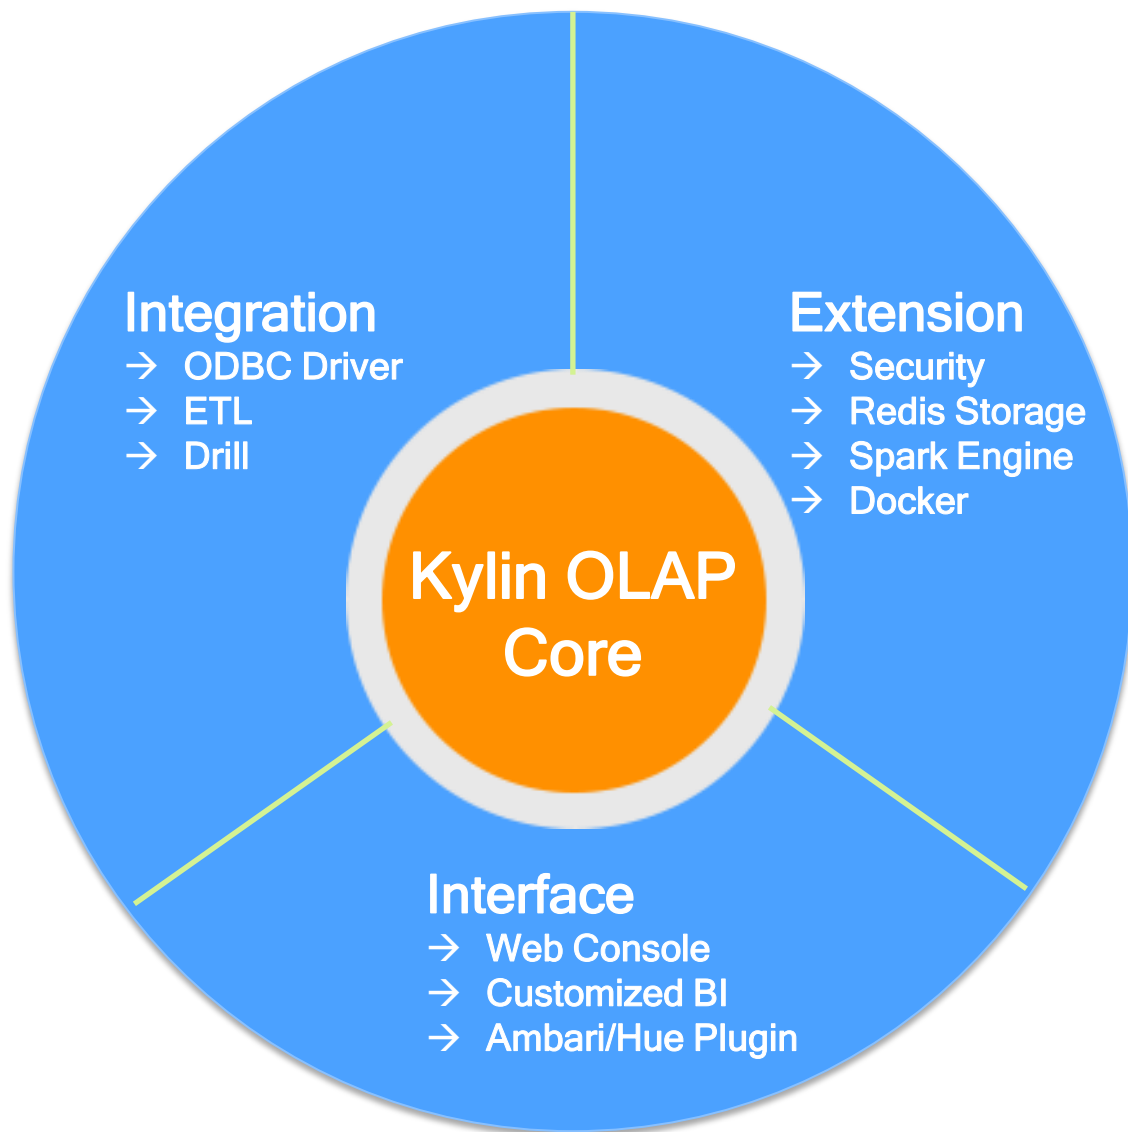
Agenda



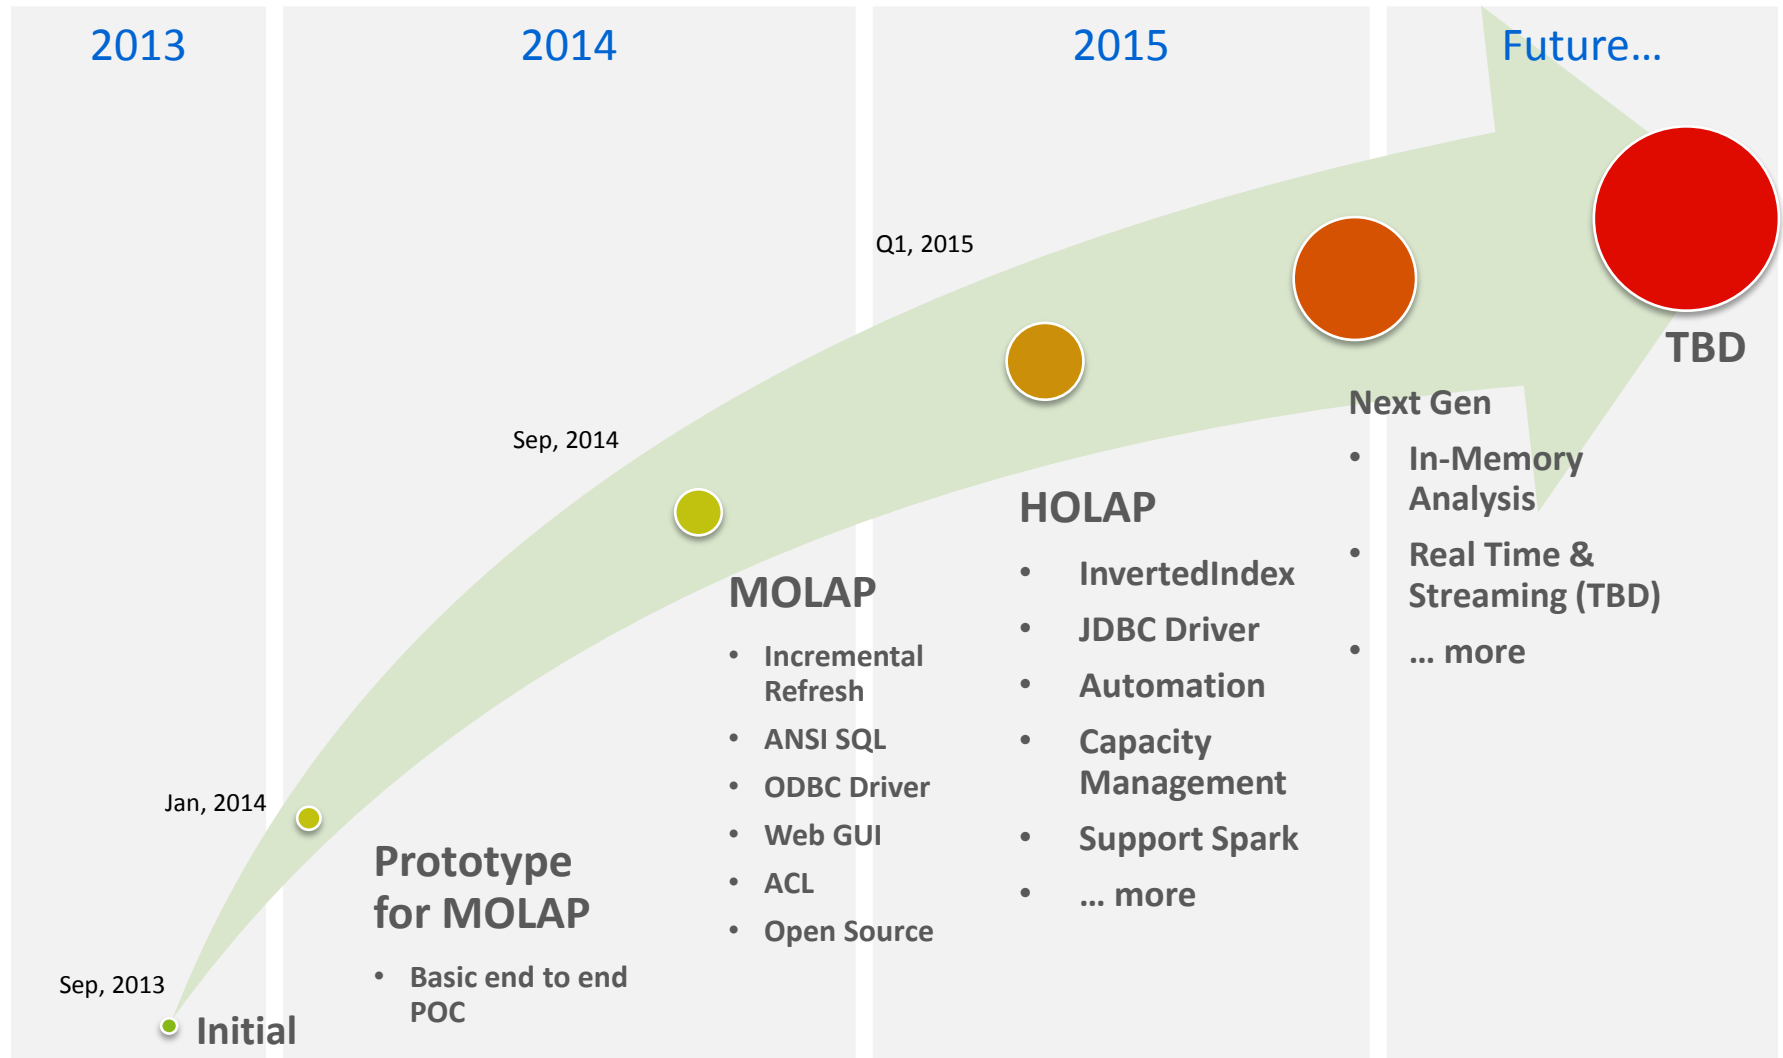
- What's Apache Kylin?
- Feature & Tech Highlights
- Performance
- Open Source & Roadmap
- Q & A

Kylin Ecosystem

- **Kylin Core**
 - Fundamental framework of Kylin OLAP Engine
- **Extension**
 - Plugins to support for additional functions and features
- **Integration**
 - Lifecycle Management
Support to integrate with other applications
- **Interface**
 - Allows for third party users to build more features via user-interface atop Kylin core
- **Driver**
 - ODBC and JDBC Drivers



Kylin Evolution Roadmap



Open Source

- Kylin Site:
 - <http://kylin.io>
- Twitter:
 - [@ApacheKylin](https://twitter.com/ApacheKylin)
- Source Code Repo:
 - <https://github.com/KylinOLAP>
- Google Group:
 - [Kylin OLAP](https://groups.google.com/forum/#!forum/kylin)



The screenshot shows the Apache Kylin website. At the top, there is a red banner with the text "Fork me on GitHub". Below the banner, the Apache Kylin logo is displayed, followed by the text "Apache Kylin" and "Extreme OLAP Engine for Big Data". A "DOWNLOAD" button is visible. The navigation menu includes "HOME", "DOCS", "COMMUNITY", "GITHUB", and "ABOUT". The main content area is titled "APACHE KYLIN OVERVIEW" and contains the following text:

Kylin has been accepted as Apache Incubator Project on Nov 25, 2014.

Apache Kylin is an open source Distributed Analytics Engine from eBay Inc. that provides SQL interface and multi-dimensional analysis (OLAP) on Hadoop supporting extremely large datasets

The diagram illustrates the Kylin architecture. It shows a central stack of components: REST Server, Query Engine, Routing, Metadata, and Cube Build Engine. The REST Server interacts with a "Third Party App (Web App, Mobile)" via a "REST API" and with "SQL-Based Tool (BI tools: Tableau...)" via "JDBC/ODBC". The Query Engine and Routing components interact with "Hadoop Hive" via "Mid Latency-Minutes" and "Star Schema Data". The Cube Build Engine interacts with "HBase as Storage" via "Key Value Data". The OLAP Cube is shown as a separate component that interacts with the central stack via "Low Latency-Seconds". A box on the right lists the "Online Analysis Data Flow" components: "Offline Data Flow", "Only SQL for End User", and "OLAP Cube is transparent to users".



Apache Kylin 北京线下交流会

- 时间:

- 2014-12-14 6:30 PM - 9:00 PM

- 地点:

- 3W咖啡



Thanks



<http://kylin.io>
lukehan@apache.org

