Hortonworks

HBase:1.0 and Beyond

Ted Yu
yuzhihong@gmail.com

About myself

- **Graduated from TsingHua University**
- **Have been working on Hbase for over four years**
- **Have been Hbase PMC member since June 2011**
- **Senior MTS at Hortonworks**

Agenda

- HBase 1.0
- HydraBase:cross DC high availability
- Local Index support in Phoenix
- Per column family flush
- Q & A

Major Changes for 1.0

- Stability: Co-locate hbase:meta with Master
- Simplify, Improve region assignment reliability: Fewer components involved
- Master embeds a RegionServer, hosting only system tables
- Backup masters can be configured to host user tables
- Plumbing is all there, **OFF** by default

- http://issues.apache.org/jira/browse/HBASE-10569

Major Changes for 1.0 (contd)

- Availability: Region Replicas
- Multiple RegionServers host a Region

① One is “primary”, others are “replicas”

② Only primary accepts writes

- Baby step toward quorum reads, writes
- Plumbing is all there, **OFF** by default

- http://issues.apache.org/jira/browse/HBASE-10070
- http://issues.apache.org/jira/browse/HBASE-11183
- http://www.slideshare.net/HBaseCon/features-session-1

Major Changes for 1.0 (contd)

- Usability: Client API changes
- Improved self-consistency
- Simpler semantics
- @InterfaceAudience annotations

- http://s.apache.org/hbase-1.0-api
- https://github.com/ndimiduk/hbase-1.0-api-examples

Client API usage sample

```
Connection conn =
    ConnectionFactory.createConnection(job.getConfiguration());
try {
  UserProvider userProvider =
    UserProvider.instantiate(job.getConfiguration());
  TokenUtil.addTokenForJob(conn, userProvider.getCurrent(), job);
} finally {
  conn.close();
}
```

http://issues.apache.org/jira/browse/HBASE-12493

Major Changes for 1.0 (contd)

- Online config change: ported from 89-fb HBASE-12147
- Automatic tuning of global MemStore and BlockCache sizes
- BucketCache easier to configure
- Pluggable replication endpoint
- Greatly expanded hbase.apache.org/book.html
- Combining mvcc/seqid
- Sundry security, tags, labels improvements

Online / Wire Compatibility

- Direct migration from 0.94 supported

① Similar to upgrade from 0.94 to 0.96: requires downtime

② Not tested yet, will be before release

- RPC is backward-compatible to 0.96

① Enabled mixing clients and servers across versions

- Rolling upgrade "out of the box" from 0.98

① 0.96 cannot read HFileV3, the new default

Client Application Compatibility

- API is backward compatible to 0.96

① No code change required

② You'll start getting new deprecation warnings

- ABI is **NOT** backward compatible

① Cannot drop current application jars onto new runtime

② Recompile your application vs. 1.0 jars

③ similar to 0.96 to 0.98 upgrade

Hadoop / Java Versions

- Hadoop 1.x is NOT supported

① you'll enjoy the performance benefits

- Hadoop 2.x only

① Most thoroughly tested on 2.4.x, 2.5.x

② less thoroughly tested on 2.2.x, 2.3.x

- JDK 6 is NOT supported!
- JDK 7 is the target runtime

- https://hbase.apache.org/book/ configuration.html#hadoop

HydraBase

- Goal: 4 9’s of availability in steady state
- No data loss at cluster level failures
- All failures should be quick to recover from
- distributed consensus shouldn’t affect write throughput
- each region will be hosted by a set of Region Servers

HydraBase: Replication Protocol

- There will be only one leader amongst the set of replicas
- Leader serves all the read and write requests to the client
- The election of the leader will be done using the RAFT protocol
- Each replica will have its own Write Ahead Log, stored locally
- Writes will be replicated synchronously by the leader to the replica set

HydraBase: RMAP

- RMap contains the quorum configuration information for each Region
- Based on the network latency to the client, each Data Center will have a rank number
- DC with the lowest network latency to the client, will have the highest rank
- Qualified quorum member with higher DC ranking is able to take over the leadership
- Replica with higher rank (DC-rank + machine-rank) will have a lower election timeout

HydraBase: Types of Replicas

- **ACTIVE:** performs all the LSM operations in the RegionServers. This includes flushes and compactions. By default, the current leader is always ACTIVE

- **ACTIVE-WITNESS:** has a memstore associated with it, but is not performing any LSM operations
- There can be one or more active-witness replicas per HDFS cluster

- **SHADOW-WITNESS:** one who is only participating in the replication via the protocol

HydraBase: Multi Cluster Deployment Setup

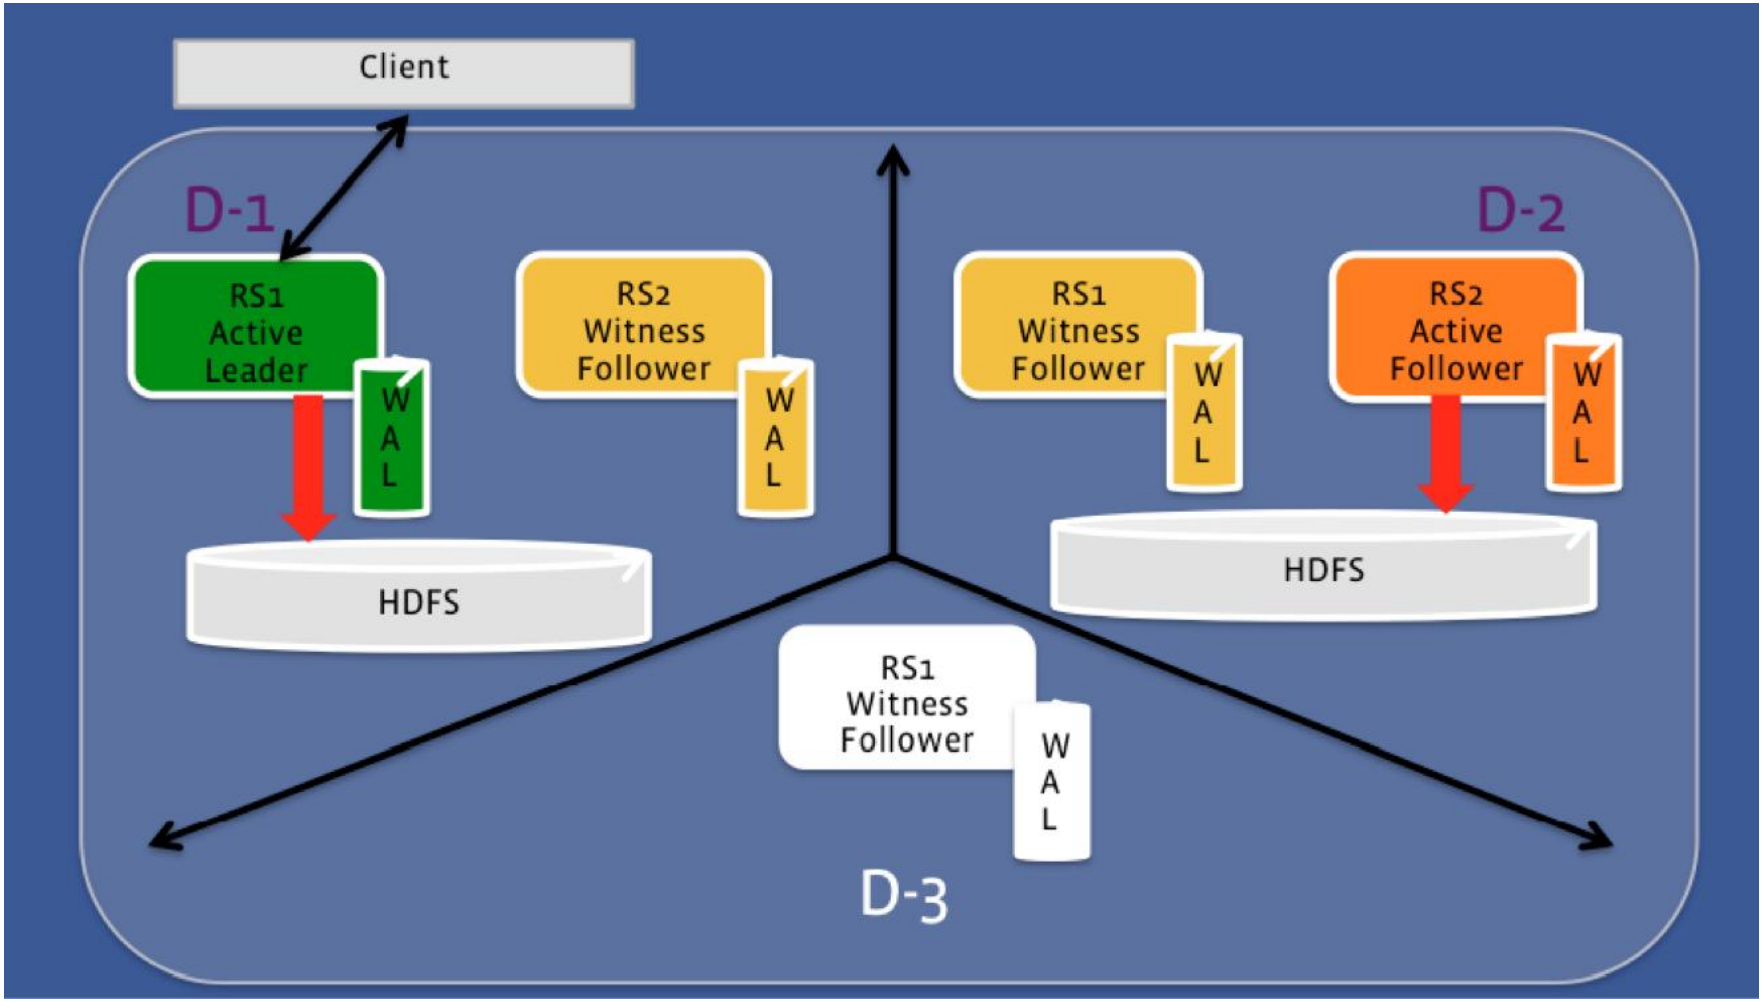

HydraBase: Single Cluster Deployment Setup

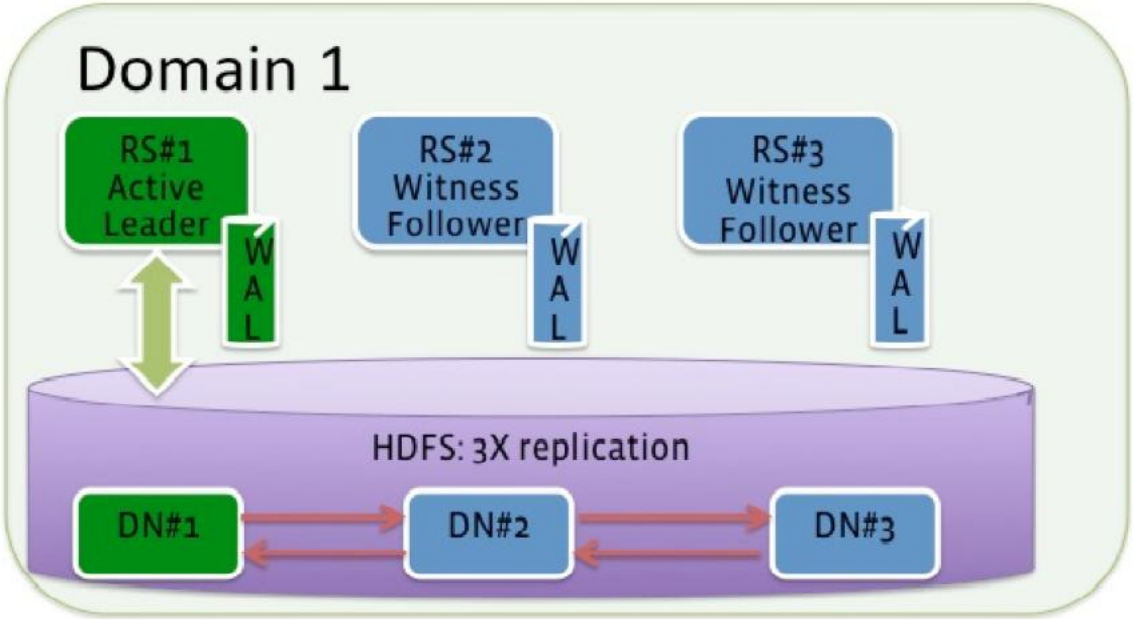

Phoenix Local Index Architecture

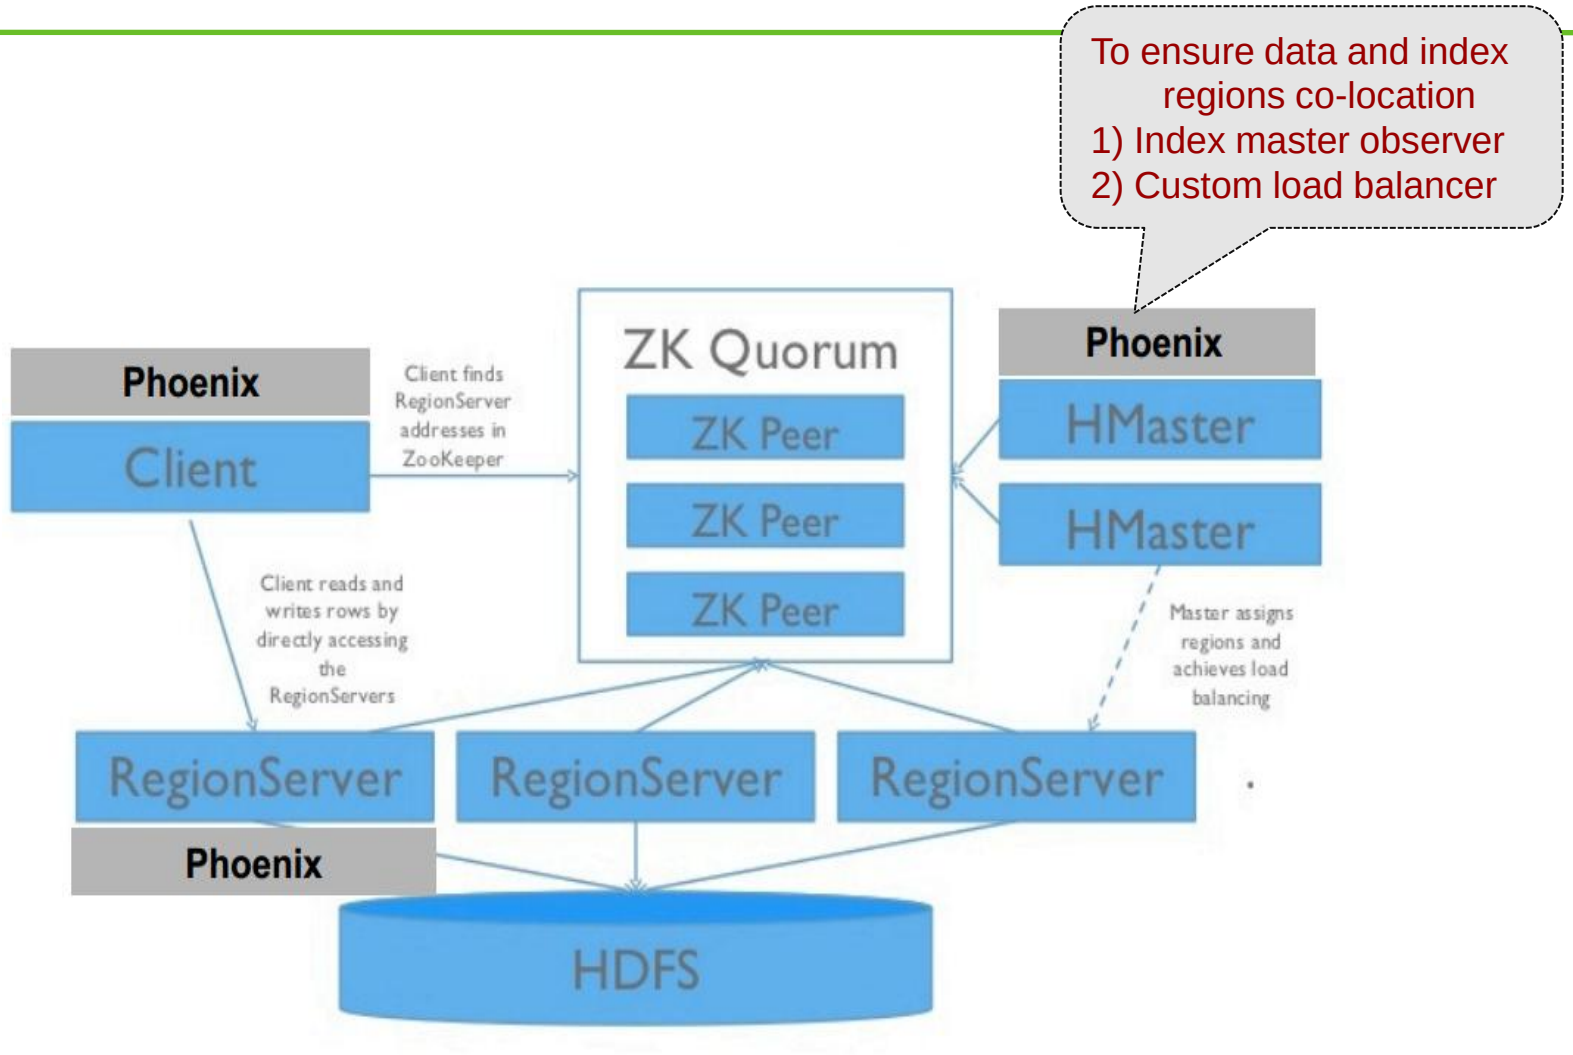

Regions Co-locate

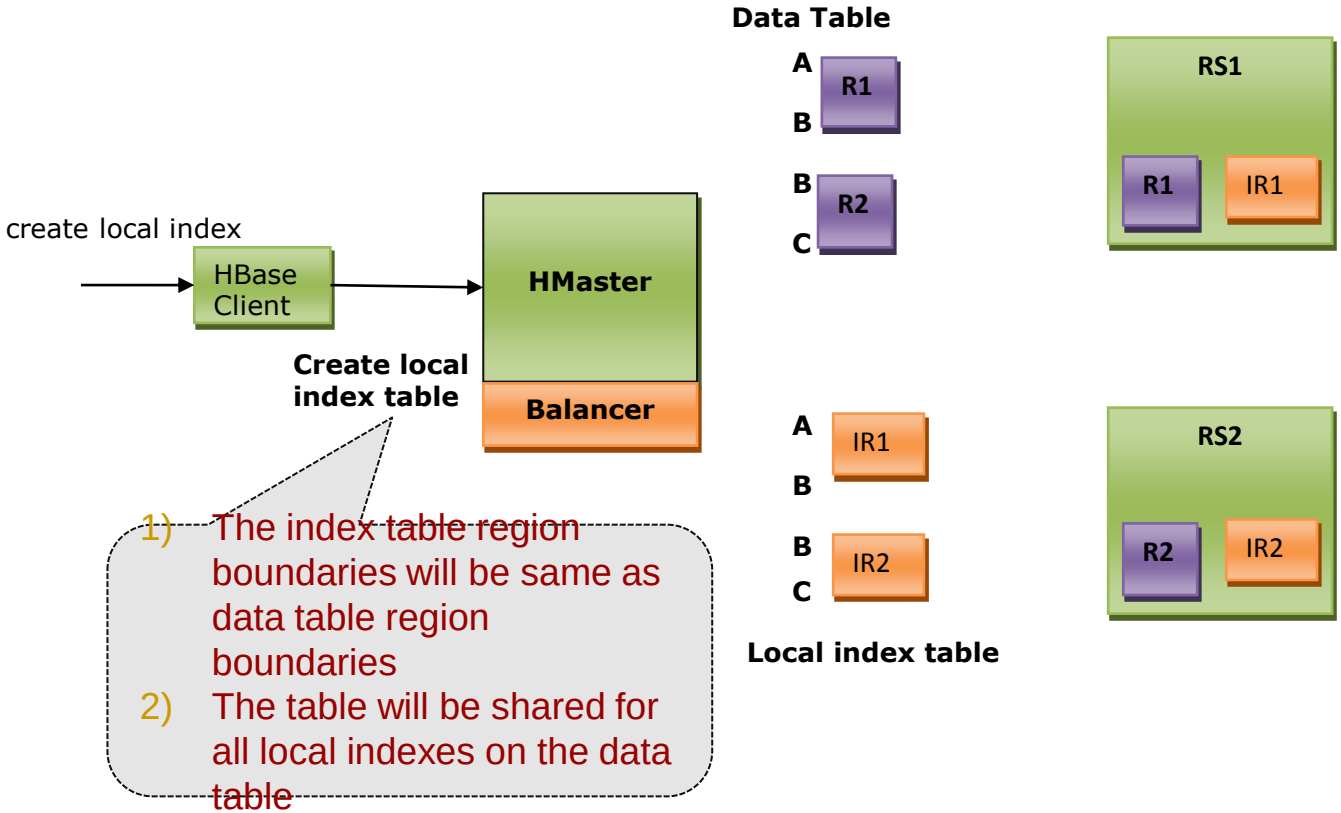

Write path

- Data table DT with columns pk,c1,c2
- Create local index LIDX on DT(c1)
- Local index table -> _LOCAL_IDX_DT

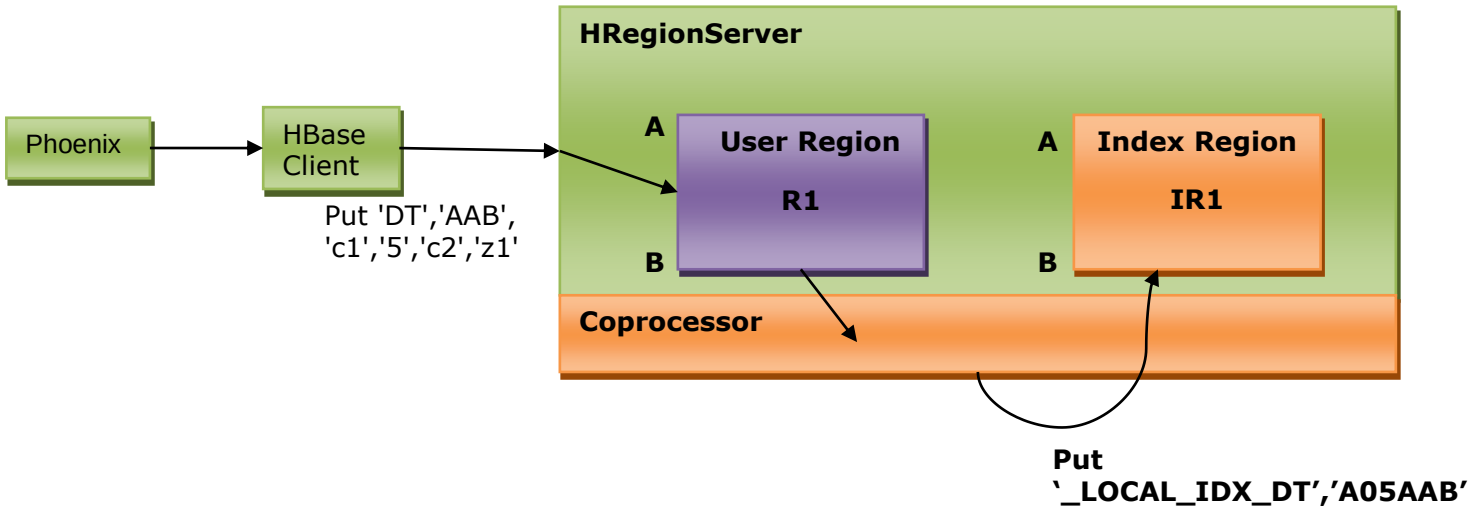

Read path

- **Select c1, c2 from DT where c1 = 5**

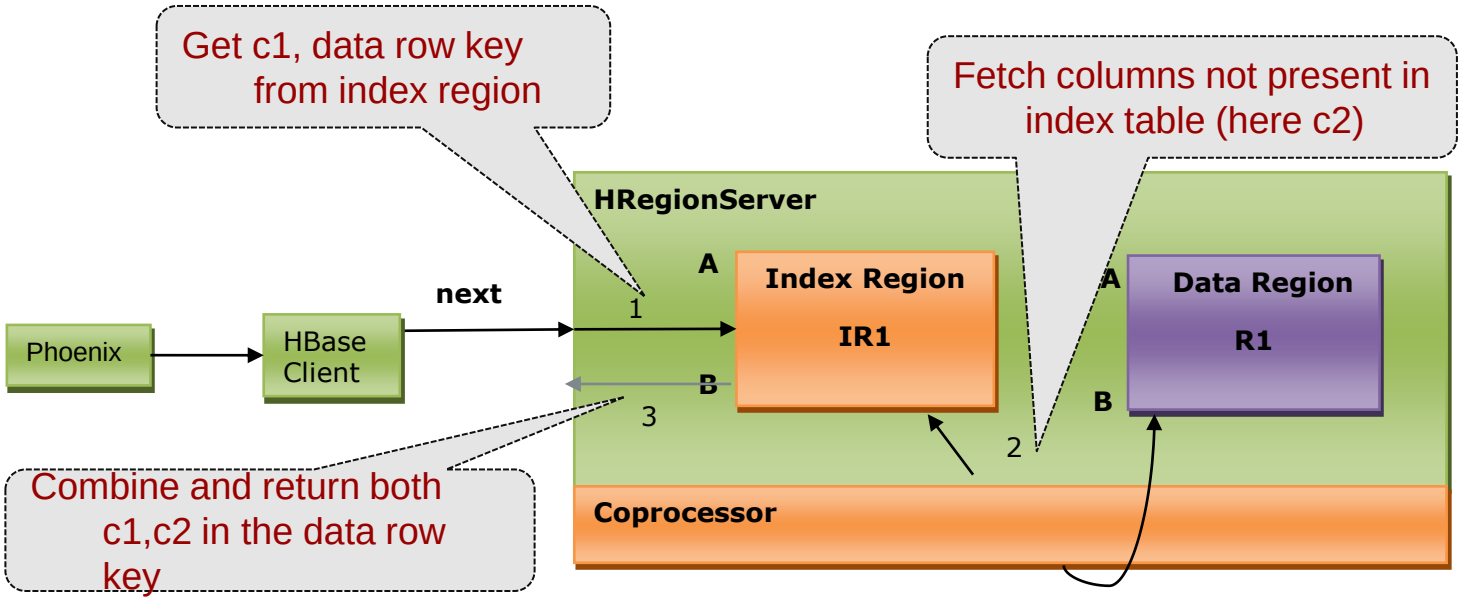

Local Index Performance

- **Created table pre split (30 regions)**
- **Number of indexes: 4**
- **Data size : 500MB**

① **No index : 4955 sec**

② **Local mutable indexes : 1152**

③ **Global mutable indexes : 1679 sec**

Multi-WAL HBASE-5699

- **multiwal-1_1_to_200_threads**

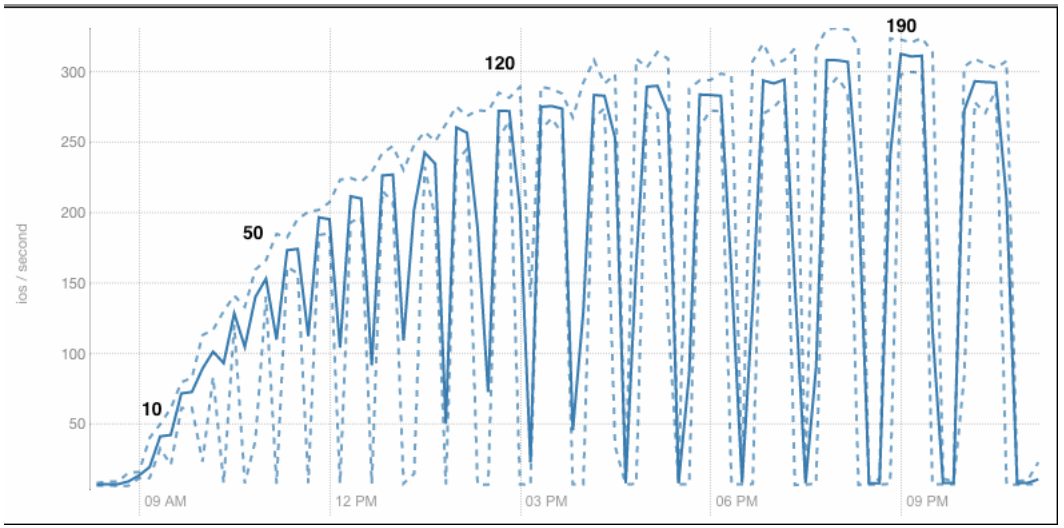

- **multiwal-2-10_to_400_threads**

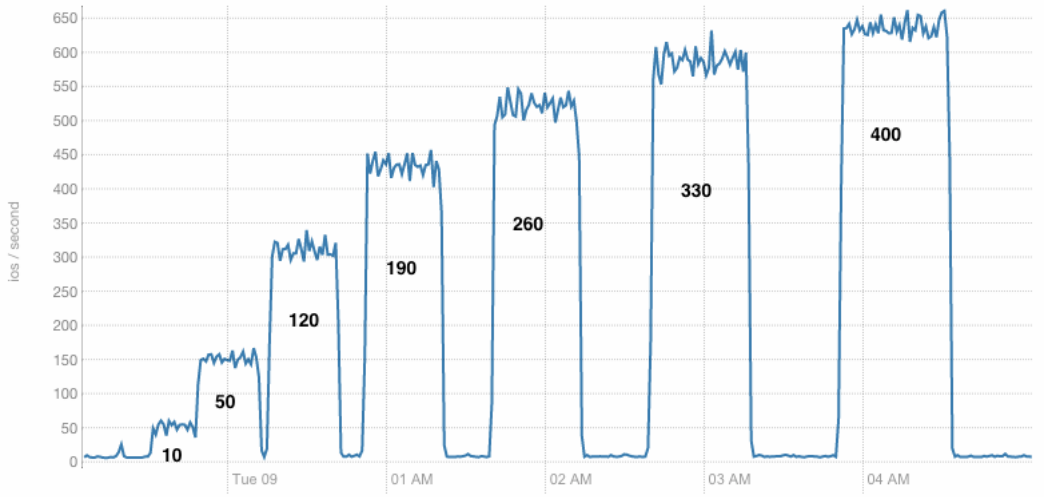

Per column family flush

- HBASE-10201 Port from 0.89-fb branch
- Reduces write amplification by 10%
- Lower bound on flush size per column family
- FlushPolicy controls whether all stores are flushed
- Make use of sequence Id of each Store

Per column family flush – I/O reduction

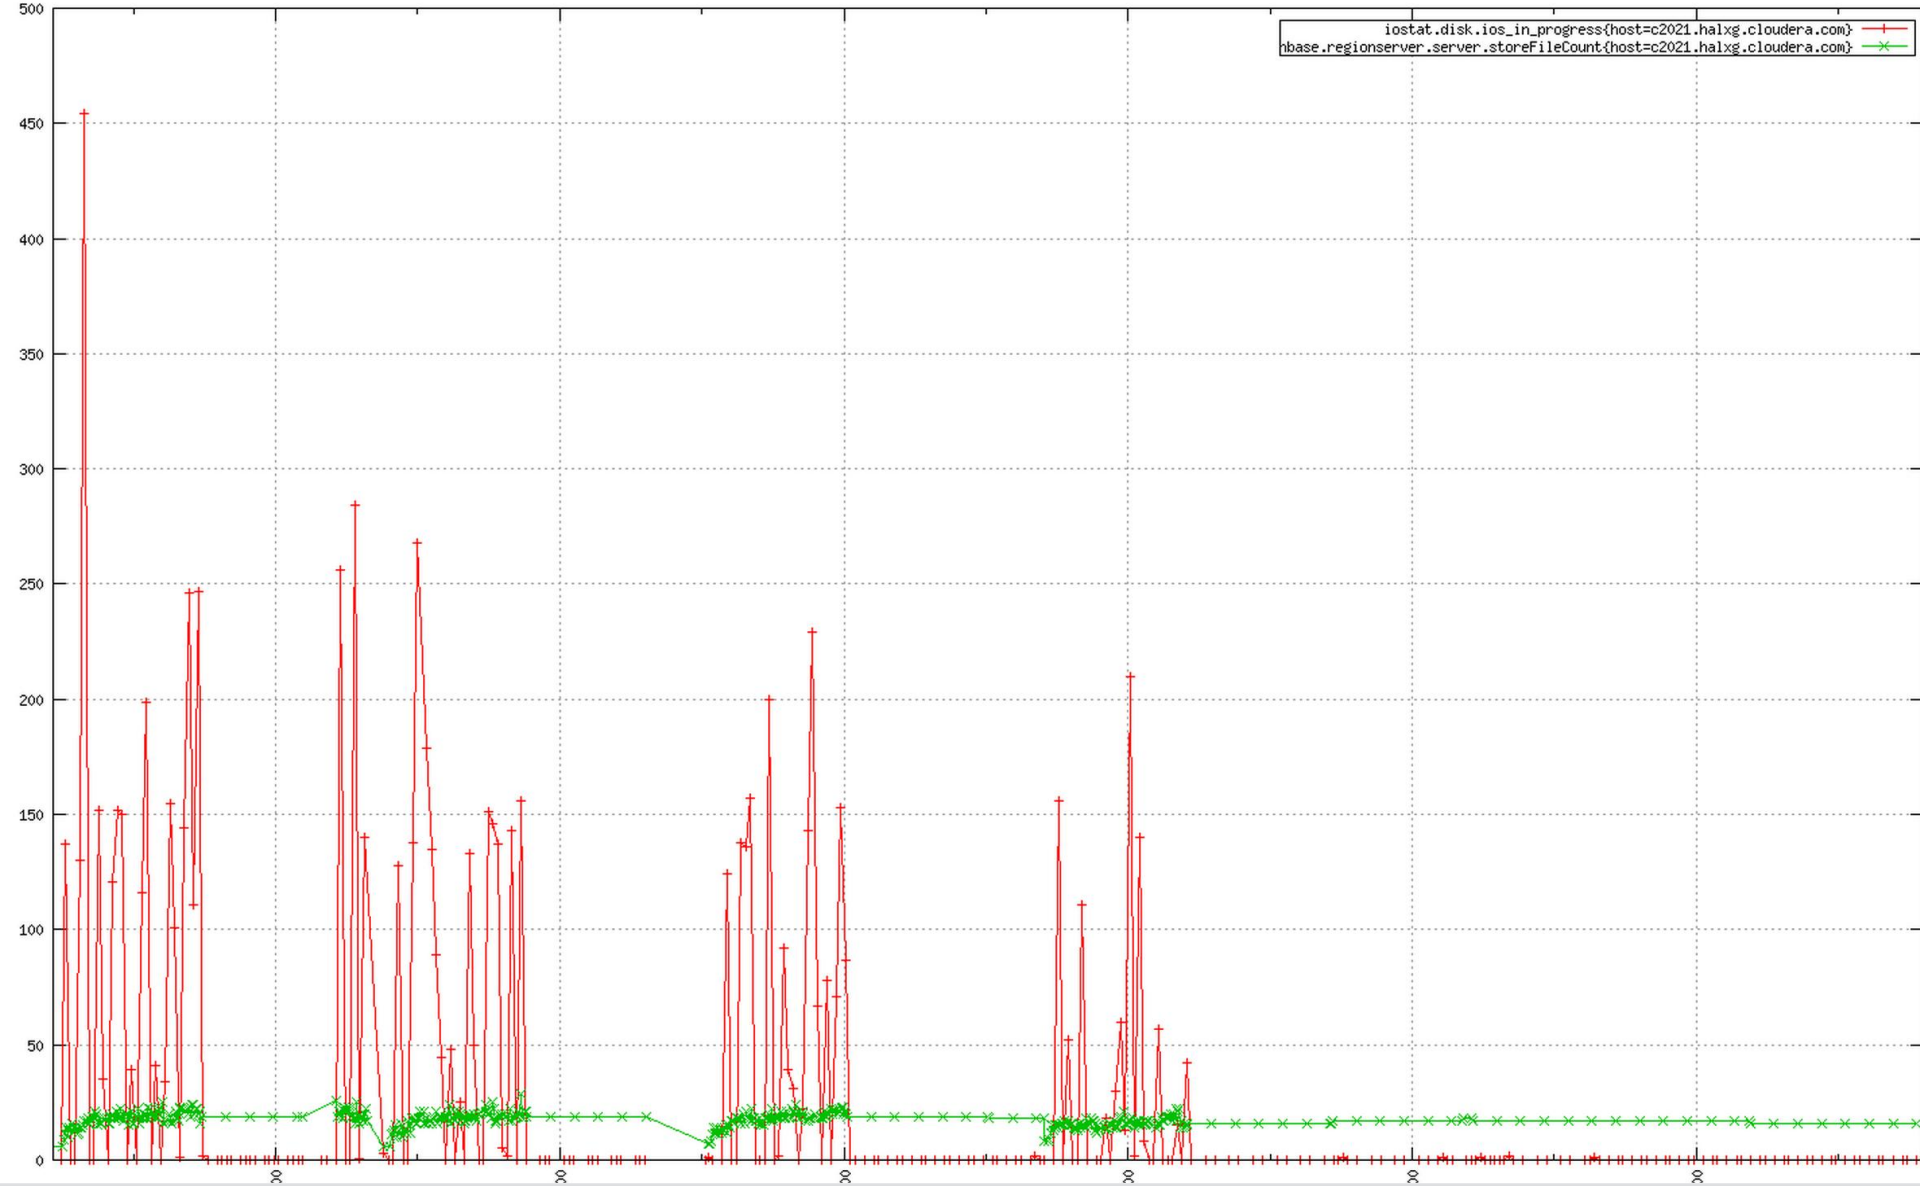

Colors

Primary Colors

Hortonworks Green
R- 105
G- 190
B- 40

Hortonworks Black
R- 30
G- 30
B- 30

Secondary Colors

Hortonworks Orange
R- 225
G- 112
B- 0

Hortonworks Dusty Blue
R- 68
G- 105
B- 125

Hortonworks Gray
R- 129
G- 138
B- 143

For PPT Only

Darker Text Gray
R- 127
G- 127
B- 127

Simple Slide: Arial, 36pt, Left Justified

- **Bulleted Body Text:**
  Arial, 18pt, Bold
  - Second Level Sub Bullet: Arial 16pt
    - Third Level Sub Bullet: Arial 14 pt
- **Bulleted Body Text:**
  Arial, 18pt, Bold
  - Second Level Sub Bullet: Arial 16pt
    - Third Level Sub Bullet: Arial 14 pt
- **Bulleted Body Text:**
  Arial, 18pt, Bold
  - Second Level Sub Bullet: Arial 16pt
    - Third Level Sub Bullet: Arial 14 pt

Hortonworks

Transition Slide: Arial 54pt

Transition Slide Sub Title: Arial 28pt

- **Bulleted Body Text:**
  Arial, 18pt, Bold
  - Second Level Sub Bullet: Arial 16pt
    - Third Level Sub Bullet: Arial 14 pt
- **Bulleted Body Text:**
  Arial, 18pt, Bold
  - Second Level Sub Bullet: Arial 16pt
    - Third Level Sub Bullet: Arial 14 pt
- **Bulleted Body Text:**
  Arial, 18pt, Bold
  - Second Level Sub Bullet: Arial 16pt
    - Third Level Sub Bullet: Arial 14 pt

Closing Slide: Arial, 54pt

Closing Sub Title: Arial, 28pt